

Tensor Cache: Eviction-conditioned Associative Memory for Transformers

Kabir Swain¹ Sijie Han² Daniel Karl I. Weidele³ Mauro Martino³ Antonio Torralba¹

Abstract

Autoregressive Transformer KV caches grow linearly with context length; sliding-window caching bounds memory but discards evicted tokens entirely, so relevant evidence outside the window becomes inaccessible. We introduce *Tensor Cache*, a two-level cache that pairs sliding-window softmax attention as a first-level cache (L1) with a fixed-size outer-product fast-weight memory as a second-level cache (L2) fed by KV pairs evicted from the window. Recent tokens remain in exact local attention; evicted pairs are compressed into a per-layer matrix A and read by future queries through a single matrix multiplication, exploiting the linear-attention identity $q_t(k_i \otimes v_i) = \langle q_t, k_i \rangle v_i$. A learned scalar gate fuses the L1 and L2 outputs, and per-head decay and write-rate parameters are trained end-to-end. The outer-product memory and the read identity are well-known; our contribution is their use as an L2 cache fed exclusively by sliding-window evictions, plus identifying that the common chunked-mean training shortcut $A \leftarrow \lambda A + \eta(\bar{k} \otimes \bar{v})$ silently introduces $C^2 - C$ spurious cross-token outer products per chunk, and closing the gap with a parallel weighted-sum scan equivalent to per-token writes within float32 epsilon. Across systems scaling, controlled associative recall, long-context language modeling, and memory-capacity diagnostics, Tensor Cache improves the memory-quality frontier over bounded-state baselines.




Figure 1. Tensor Cache. Each layer keeps a local KV ring buffer (L1) and a fixed-size memory A (L2). On eviction, (k, v) is written into A via an outer-product update; queries read both paths and fuse via a learned gate.

1. Introduction

Autoregressive Transformer inference caches per-layer keys and values (KV) so the prefix is not recomputed at every step (Vaswani et al., 2017; Shazeer, 2019), but retained KV state grows linearly with context length, depth, head dimension, and serving concurrency (Kwon et al., 2023). Sliding-window caching bounds this cost by keeping only the most recent W tokens (Xiao et al., 2023), at the price of *hard forgetting*: once a token leaves the window, its key and value are removed from attention entirely.

We introduce *Tensor Cache* (TC), a bounded-state mechanism that turns KV eviction from a deletion event into an associative-memory write. Each layer keeps exact softmax attention over the recent window (a first-level cache, L1), and when the ring buffer overwrites an old entry, the evicted pair (k_w, v_w) is written into a fixed-size outer-product matrix A (a second-level cache, L2) of the kind used in fast-weight programmers (Schmidhuber, 1992) and modern linear-attention variants (Schlag et al., 2021; Katharopoulos et al., 2020; Beck et al., 2024; Munkhdalai et al., 2024). Future queries read this state by

$$q_t A \approx \sum_i \alpha_i \langle q_t, k_i \rangle v_i,$$

i.e. the same query-key alignment used by attention also addresses the compressed past, but without retaining the

¹Massachusetts Institute of Technology, Cambridge, MA, USA ²University of Toronto, Toronto, Canada ³IBM Research, Cambridge, MA, USA. Correspondence to: Kabir Swain <kswain@mit.edu>, Sijie Han <hs.han@mail.utoronto.ca>, Daniel Karl I. Weidele <daniel.karl@ibm.com>, Mauro Martino <mmartino@us.ibm.com>, Antonio Torralba <torralba@mit.edu>.

individual evicted entries. A learned scalar gate fuses the L1 and L2 outputs.

The outer-product memory and read identity $q_t(k_i \otimes v_i) = \langle q_t, k_i \rangle v_i$ are not new. Our contribution is to use them as an L2 cache fed exclusively by sliding-window evictions, rather than as a parallel running average over every token (mLSTM, Infini-attention) or as a wholesale replacement for attention (RetNet, Mamba) (Sun et al., 2023; Gu & Dao, 2023). We additionally identify a previously-overlooked train-time shortcut: a naive chunked-scan implementation summarizes each chunk by its mean and writes $A \leftarrow \lambda A + \eta(\bar{k} \otimes \bar{v})$, which differs from the per-token sum by $C^2 - C$ cross-token outer products $k_i \otimes v_j$ ($i \neq j$) that never appear at inference. We close this gap with a parallel weighted-sum scan that is mathematically equivalent to per-token writes for the outer rule, verified to float32 epsilon. The parallel-scan formulation itself is standard in the linear-attention literature (Beck et al., 2024; Munkhdalai et al., 2024); we make its applicability to eviction-conditioned memory explicit and quantify the cost of skipping it.

TC is not a lossless replacement for full attention. Full KV remains the exact-content baseline when affordable; in evaluations beyond the training block size, Full-KV results should be read as exact-content but positionally extrapolated (Su et al., 2021; Chen et al., 2023; Peng et al., 2024). TC targets the regime where full-KV state is too expensive but pure sliding-window decoding forgets too aggressively.

Contributions.

- A two-level cache architecture combining sliding-window softmax attention (L1) with an outer-product fast-weight memory (L2) fed exclusively by KV pairs evicted from the window. Unlike prior outer-product memories that ingest every token, our eviction-conditioned write makes the matrix specifically a second-level store for what local attention can no longer reach.
- Identification and closure of a chunked-mean training shortcut. The shortcut introduces $C^2 - C$ spurious cross-token outer products per chunk; the parallel weighted-sum scan removes them at the same compute cost (verified within 10^{-7} relative error).
- Empirical evaluation across systems scaling, real-text long-context language modeling, controlled associative recall, raw memory capacity, and ablations. Tensor Cache attains 100% matched-gap synthetic recall while using 72–84% less inference state than Full KV, and on OpenWebText long-context language modeling it attains the lowest mean NLL of any tested method—including Full KV, Window KV, StreamingLLM, and Infini-attention—at every evaluation context length

from 1024 to 32,768 tokens with the single exception of $L = 2048$, where it is statistically indistinguishable from Infini-attention; at $32\times$ the trained context Tensor Cache reaches NLL 5.14 versus 6.00 for Full KV at $2.4\times$ greater peak GPU memory.

2. Related Work

KV cache systems and eviction. KV-cache footprint scales with context length and serving concurrency. Systems work such as PagedAttention (Kwon et al., 2023) and CacheGen (Liu et al., 2023b) improves allocation, compression, or movement, but the logical retained state still scales with the amount of context kept. A complementary line reduces KV size by retaining only selected tokens: sliding-window attention (Longformer (Beltagy et al., 2020), StreamingLLM with attention sinks (Xiao et al., 2023)) and importance-based selection (H₂O, SnapKV, CAOTE (Zhang et al., 2023; Li et al., 2024; Goel et al., 2025)). These methods decide which discrete entries to keep. Tensor Cache instead accepts FIFO eviction and converts the evicted pair into a compact associative-memory update.

Memory-augmented and recurrent Transformers.

Transformer-XL reuses hidden states across segments (Dai et al., 2019), Compressive Transformers compress older activations (Rae et al., 2019), Memorizing Transformers use external kNN memory (Wu et al., 2022), and Recurrent Memory Transformer carries information through memory tokens (Bulatov et al., 2022). Most relevant to bounded-state decoding, LESS combines sparse KV retention with a low-rank recurrent cache (Dong et al., 2024), and Infini-attention combines masked local attention with a bounded compressive memory updated every token through a normalized outer-product write (Munkhdalai et al., 2024). Tensor Cache differs in the semantics of the write: a key–value pair is first used in exact local softmax attention, and only when the ring buffer overwrites it is the pair written into the auxiliary state.

Fast weights, outer-product memories, and linear attention.

The outer-product fast-weight memory traces back to Schmidhuber’s fast-weight programmer (Schmidhuber, 1992); Ba et al. (2016) revisited fast weights as short-term memory. Widrow–Hoff (Widrow & Hoff, 1960) introduced the error-correcting principle behind the delta rule. Schlag et al. (2021) reframed linear attention (Katharopoulos et al., 2020) as a fast-weight programmer with accumulated $k \otimes v$ writes, formalizing the read identity $q(k \otimes v) = \langle q, k \rangle v$ that we exploit. Yang et al. (2024) developed parallel delta-rule matrix-state models, and mLSTM (Beck et al., 2024) revisits matrix-valued recurrent memory with input/forget gating and a normalization vector. Tensor Cache borrows the matrix primitive and the read identity from this line of

work; its architectural distinction is to use the matrix as a *second-level cache* fed by sliding-window evictions, rather than as an attention replacement (mLSTM, RetNet (Sun et al., 2023), Mamba (Gu & Dao, 2023)) or as a parallel running average over every token (Infini-attention). We do not require Infini’s normalization vector or feature map, instead using a learned exponential decay.

Training-side considerations. mLSTM and Infini-attention’s segment scan preserve per-token write semantics during chunked training. We observed that a chunk-mean shortcut $A \leftarrow \lambda A + \eta(\bar{k} \otimes \bar{v})$ is a tempting but lossy approximation that introduces $C^2 - C$ spurious cross-token outer products per chunk and degrades retrieval when an associative episode fits within a single chunk. Replacing the chunk-mean write with the standard parallel weighted-sum scan—mathematically equivalent to per-token writes for the outer rule—closes the gap. The scan itself is standard linear-recurrence material; we make its applicability to eviction-conditioned outer-product memories explicit and quantify the cost of skipping it.

Other long-context methods. Ring Attention (Liu et al., 2023a) distributes blockwise attention; LongNet (Ding et al., 2023) uses dilated attention; FlashAttention-2 (Dao, 2023) improves throughput. These are largely orthogonal: Tensor Cache augments a sliding-window attention stack with a fixed-size associative memory path for evicted KV entries and could be combined with such systems-level optimizations.

3. Architecture

Tensor Cache (TC) augments each Transformer layer with three components: a fixed-capacity KV ring buffer for exact local attention (L1), a fixed-size associative memory state for evicted information (L2), and a learned gate that fuses the two.

3.1. Setup and local KV window

Let B be batch size, H the number of heads, D the head dimension. Each layer produces $Q, K, V \in \mathbb{R}^{B \times H \times T \times D}$ from a single input projection, optionally followed by RoPE on Q, K . The local window

$$K_{\text{win}}, V_{\text{win}} \in \mathbb{R}^{B \times H \times W \times D} \quad (1)$$

holds the last W key/value pairs and is used for exact causal attention $y_t^{\text{local}} = \text{Attn}(q_t, K_{\text{win}}, V_{\text{win}})$. RoPE is applied before keys enter the buffer, so the keys later written to memory are post-RoPE.

3.2. Eviction-conditioned memory state

Each layer keeps a fast-weight matrix $A \in \mathbb{R}^{B \times H \times D \times D}$, initialized to $A_0 = \mathbf{0}$. The KV cache is a FIFO ring buffer; when slot pos is overwritten, the displaced pair $(k_{\text{old}}, v_{\text{old}})$ becomes the TC write source (k_w, v_w) (*write-on-evict*). This is the architectural distinction between TC and prior outer-product memories such as mLSTM (Beck et al., 2024) and Infini-attention (Munkhdalai et al., 2024), which write every token unconditionally; TC writes only what the local window discards.

3.3. Memory read and write

The current query reads $r_t = q_t A \in \mathbb{R}^{B \times H \times D}$. After accumulated outer-product writes $A \approx \sum_i \alpha_i (k_i \otimes v_i)$, the read expands as

$$r_t = q_t A \approx \sum_i \alpha_i \langle q_t, k_i \rangle v_i, \quad (2)$$

using the standard linear-attention identity $q(k \otimes v) = \langle q, k \rangle v$ (Schmidhuber, 1992; Schlag et al., 2021). The default *outer rule* updates

$$A \leftarrow \lambda A + \eta (k_w \otimes v_w), \quad (3)$$

where λ and η are sigmoid-parameterized per-head scalars, $\lambda_h = \sigma(\theta_h^\lambda)$ and $\eta_h = \sigma(\theta_h^\eta)$ for $h = 1, \dots, H$. We also implement a *delta-rule* alternative

$$A \leftarrow \lambda A + \eta k_w \otimes (v_w - k_w A), \quad (4)$$

which writes only the residual against the memory’s current prediction (Widrow & Hoff, 1960; Schlag et al., 2021). We use the outer rule by default because it admits the exact parallel scan below; the delta rule requires a chunk-start residual approximation. We ablate both (Section 4.6).

The memory read is merged across heads, projected by a separate W_{tc} , and fused with local attention by a learned scalar gate g :

$$y_t = y_t^{\text{local}} + \sigma(g) m_t, \quad m_t = \text{MergeHeads}(r_t) W_{\text{tc}}. \quad (5)$$

3.4. Parallel chunked-scan training

Streaming inference contributes one rank-one outer product per evicted token. A natural batched-training shortcut is to summarize each chunk by its mean and write a single per-chunk outer product:

$$A \leftarrow \lambda A + \eta (\bar{k} \otimes \bar{v}), \quad \bar{k} = \frac{1}{C} \sum_t k_t, \quad \bar{v} = \frac{1}{C} \sum_t v_t. \quad (6)$$

Writing out the means exposes the problem:

$$\bar{k} \otimes \bar{v} = \frac{1}{C^2} \sum_{i,j=1}^C k_i \otimes v_j \neq \frac{1}{C} \sum_{t=1}^C k_t \otimes v_t. \quad (7)$$




Figure 2. One streaming step of Tensor Cache, shown left to right. **(1) Local window full:** the local KV ring buffer holds the most recent W key/value pairs as a new pair (k_t, v_t) arrives. **(2) Evict oldest:** the displaced pair (k_{old}, v_{old}) is popped from the buffer to make room. **(3) Update Tensor Cache:** the evicted pair is written into the L2 attention memory A via the outer-product (or optional delta-rule) update $A \leftarrow \lambda A + \eta \text{update}(k_{old}, v_{old})$, compressing long-term information. **(4) Insert new token:** (k_t, v_t) takes the freshest slot of the ring buffer. **(5) Read and fuse:** the current query q_t reads both paths — local window output y_t^{local} and memory output $m_t = q_t A$ — which are combined via a learned scalar gate, $y_t = y_t^{\text{local}} + \sigma(g) m_t$.

The $C^2 - C$ cross-token terms $k_i \otimes v_j$ with $i \neq j$ do not occur during per-token streaming inference, so chunked-mean training fits a state distribution that does not match deployment. Empirically, this collapses synthetic associative recall when an entire “store–filler–query” episode fits within a chunk (Section 4.6).

We close the gap with a parallel weighted-sum scan equivalent to per-token writes. For a chunk of length C , define decay weights $w_t = \lambda^{C-1-t}$ and update

$$A \leftarrow \lambda^C A + \eta \sum_{t=0}^{C-1} w_t (k_t \otimes v_t), \quad (8)$$

which expands to the same closed form as C sequential per-token outer-rule updates. The weighted sum is a single batched einsum, so chunked compute cost is preserved while the cross-term mismatch is removed. We verify the equivalence numerically: (8) reproduces the sequential per-token recurrence within float32 epsilon (relative error $< 10^{-7}$). The same scan principle is used by mLSTM (Beck et al., 2024) and Infini-attention’s segment scan (Munkhdalai et al., 2024); we make its applicability to eviction-conditioned outer-product memory explicit and ablate the chunked-mean baseline directly. For the delta rule, $\hat{v}_t = k_t A_{t-1}$ depends on the running A and is not exactly parallelizable; we use the standard chunk-start approximation (Beck et al., 2024) that re-uses the chunk-start A in $\hat{v}_t = k_t A$ for every t in the chunk:

$$A \leftarrow \lambda^C A + \eta \sum_{t=0}^{C-1} w_t k_t \otimes (v_t - k_t A), \quad (9)$$

which is empirically tight at moderate chunk sizes (Section 4.6).

Streaming inference applies (3) only on eviction (between writes, A is unchanged), so A at step $t > W$ contains contributions only from tokens that have left the local window. In

training, the chunked scan applies (8) to *every* token in the chunk, so A accumulates writes from all tokens, including those still inside the local window. The two regimes share the read identity (2) and the per-token mathematical form, but the population of tokens written into A differs.

The full-forward path is differentiable through local attention, the TC read/write branch, the per-head λ, η , the fusion gate, and W_{tc} ; streaming prefill and generation run under `torch.no_grad()`.

We defer multi-slot routing, two-timescale memory, layer-selection, and key/value normalization variants to Appendix A.

4. Experiments

We evaluate Tensor Cache along four axes: retained-state scaling, controlled associative recall, real-text long-context quality, and raw associative-memory capacity. Long-gap recall, multistore stress tests, training efficiency, and extended ablations are in Appendix B.

4.1. Experimental Setup

Methods. We compare five methods: **Full KV** (retains all tokens); **Window KV** (keeps the most recent W); **StreamingLLM** (recent window plus attention sinks); **InfiniAttention** (bounded compressive memory updated every token via a normalized outer-product write); and **Tensor Cache** (same local window as the bounded baselines, with evicted KV entries written into a fixed-size outer-product memory via the parallel chunked scan of Section 3). Unless stated otherwise, TC uses the outer rule (Eq. (3)), chunk size $C = 32$, and write-on-evict. All methods share the same backbone, tokenizer, optimizer, precision, and training budget; bounded methods share the same local window so differences reflect compression of older context, not re-




Figure 3. Retained **inference** state (GB) versus context length, from 4K to 128K tokens. Full KV grows approximately linearly with context, rising from ~ 0.25 GB at $L = 4$ K to ~ 7.75 GB at $L = 128$ K, and crosses above the bounded-memory methods near $L \approx 12$ K. The four bounded methods — Window KV, StreamingLLM, Infini-attention, and Tensor Cache — remain approximately constant at ~ 0.75 – 0.80 GB across the full range. **Inset:** zoom into the bounded-memory region; Tensor Cache adds a small constant overhead (~ 0.04 GB above Window KV/StreamingLLM) corresponding to the per-layer associative-memory matrices.

tention of recent context, and are evaluated through their streaming inference path. Beyond the training block size, Full KV uses NTK-aware RoPE scaling; comparisons in that regime are against the extrapolated Full-KV baseline, not exact attention.

Datasets. Real-text training: OpenWebText (Gokaslan & Cohen, 2019) (large-scale, §4.4) and Shakespeare (converged-regime replication, §4.4); seeds {1337, 2027, 3037}. Synthetic associative-recall tasks isolate beyond-window retrieval; the main task uses matched gaps $g \in \{24, 36, 48\}$ with $W = 12$, and we additionally sweep evaluation filler lengths for OOD gap generalization.

4.2. Bounded-State Systems Scaling

Retained inference state grows linearly for Full KV and stays approximately constant for all bounded methods at context lengths from 4K to 128K (Figure 3); Tensor Cache adds a small constant overhead from the per-layer associative-memory matrices.

Streaming-decode throughput. Figure 4 reports decode tokens-per-second on the OpenWebText evaluation across $L \in [1, 024, 32, 768]$. Bounded methods are approximately flat; Full KV degrades with context and crosses below Window KV at $L = 32,768$. Tensor Cache pays a $\sim 30\%$




Figure 4. **Streaming-decode throughput vs context length** on the OpenWebText long-context evaluation (130M params, $W = 512$; median decode tokens-per-second over four eval seeds). Full KV starts highest (~ 127 tok/s at $L = 1$ K) but degrades with context and crosses below Window KV at $L = 32$ K (~ 106 vs. ~ 115 tok/s). All bounded methods remain approximately flat across the full range: Window KV (~ 115 tok/s), StreamingLLM (~ 114 tok/s), Tensor Cache (~ 80 tok/s), Infini-attention (~ 74 tok/s). Tensor Cache pays a $\sim 30\%$ throughput overhead vs. Window KV/StreamingLLM corresponding to the L2-cache read and update, and is consistently faster than Infini-attention at every L tested.

throughput overhead vs. Window KV/StreamingLLM (the L2 read/update) and is faster than Infini-attention at every length tested.

4.3. Synthetic Associative Recall

Task and setup. Each sequence contains $E = 6$ episodes of the form

$$[\text{STORE}, k, v, \text{GAP}, f_1, \dots, f_g, \text{QUERY}, k, \text{ANSWER}, v]$$

over 16 keys, values, and fillers, with only the answer token supervised (chance $1/16 = 6.25\%$). The filler exceeds the sliding KV window, so windowed methods cannot solve the task without auxiliary memory. We train at matched gaps $g \in \{24, 36, 48\}$ with a 4-layer, 4-head, 128-dim Transformer, RoPE, AdamW at 10^{-3} , batch 32, bfloat16, 300 steps, 3 seeds. Window KV uses $W = 12$; StreamingLLM uses 4 sinks plus a 12-token window; InfiniAttention and TC share the same local window. Streaming state is the peak retained inference state per sequence, including the KV window and any auxiliary memory tensors.

Results. Table 1 reports answer accuracy and per-sequence streaming state at each matched gap. Tensor Cache matches Full KV at all gaps while reducing retained inference state from 398/542/686 kB to a constant 112 kB per sequence—savings of 71.9%, 79.3%, and 83.7%, respectively. Window KV and StreamingLLM remain near chance ($\approx 6\%$) once the query lies outside the 12-token local window. InfiniAttention degrades progressively as the gap

increases, falling from 39.6% at gap 24 to 11.9% at gap 48, despite using nearly identical compact state (114 kB).

OOD calibration. To probe gap generalization, we sweep evaluation filler length $F_{\text{eval}} \in \{12, \dots, 192\}$ at training gap $g = 24$. At in-window OOD ($F_{\text{eval}} = 12$), Infini-attention attains 4.6% accuracy at NLL 14.3 — worse than uniform ($\log N \approx 2.77$), i.e. confidently incorrect retrieval. Tensor Cache attains 25.4% at NLL 6.4; with no read-time normalization, TC avoids the unbounded-magnitude crosstalk that arises on OOD queries. TC state stays at 112 kB across all F_{eval} ; Full KV grows from 254 kB at $F_{\text{eval}} = 12$ to 2,414 kB at $F_{\text{eval}} = 192$ ($22\times$ less state at the longest gap, within 5 pp of accuracy).

4.4. Long-Context Language Modeling

Setup. We train each method from scratch on OpenWebText (block size 1,024) and evaluate streaming NLL up to 32,768 tokens. Full KV uses NTK-aware RoPE scaling at evaluation; linear-attention methods need no such adjustment.

Results. Table 2 and Figure 5 show the memory–quality frontier. **Tensor Cache attains the lowest mean NLL at every evaluation L except $L = 2,048$** , where it is essentially tied with Window KV (TC 4.00 vs. 3.98, within single-seed eval-position variance). At $L = 32,768$, TC reaches NLL 5.14 — 0.86 below Full KV (6.00), 0.28 below Infini-attention (5.42), and 0.30 below Window KV (5.44) — at $2.4\times$ less peak GPU memory than Full KV. The TC margin to the next-best constant-memory baseline (Infini-attention) is 0.28 NLL at both $L = 8192$ and $L = 32,768$, consistent with the L2 cache contributing exactly where the KV window ($W = 512$) cannot.

Evaluation variance. TC’s in-window cells ($L \leq 4096$) are mean $\pm$ std over 4 eval seeds (resampling the validation start position); long-context cells and all baseline cells are single-seed (baseline checkpoints were unavailable for re-evaluation). Eval-position variance dominates in-window: TC’s NLL at $L = 1024$ ranges from 3.28 to 4.81 across seeds (std 0.66), shrinking to std 0.21 at $L = 4096$. Window KV’s $L = 2048$ value (3.98) is anomalously low relative to its neighbors (4.86 at $L = 1024$, 5.06 at $L = 4096$) and to every other method’s trajectory — likely a similar single-seed artifact. The long-context regime ($L \geq 4096$) is therefore the primary comparison: TC’s margin to every baseline there exceeds any plausible eval-noise envelope.

Replication on Shakespeare (converged regime). We additionally train all five methods from scratch on Shakespeare (300K tokens; 4-layer 128-dim, $W = 128$, block size 256, 3000 steps), with 4 eval seeds at every context




Figure 5. Long-context quality. Streaming NLL versus evaluation context length for each method. **Solid lines (foreground)** show the OpenWebText evaluation (Table 2 panel a), and **faded lines (background)** show the Shakespeare evaluation (Table 2 panel b) for visual context. The dotted vertical line marks the OpenWebText trained context length (1,024 tokens). Full KV degrades substantially past the trained length even with NTK-aware RoPE scaling; the constant-memory methods remain stable. Tensor Cache attains the lowest NLL at every evaluated context length on both datasets, with the single exception of $L = 2,048$ on OpenWebText where it ties Window KV (4.00 vs. 3.98).

length. Lower half of Table 2: **TC attains the lowest mean NLL at every L from 256 to 16,384** ($64\times$ the trained block size), with margins of 0.01–0.22 NLL over Infini-attention and 0.12–0.61 over Window-KV. The TC margin grows steadily from $L = 512$ (0.08 over Infini) to $L = 8192$ (0.22), then settles at 0.19 at $L = 16,384$ — consistent with the L2 cache compounding past the trained block size and saturating at extreme extrapolation.

4.5. Raw Associative-Memory Capacity

We isolate the storage capacity of the Tensor Cache fast-weight state independent of the Transformer. Each attention head maintains a fixed-size associative memory matrix $A \in \mathbb{R}^{D \times D}$, where D is the head dimension. We write synthetic key–value pairs directly into this memory using

$$A_t = \lambda A_{t-1} + \eta k_t v_t^\top,$$

and read with

$$r_t = k^\top A_t.$$

After each write, we query the oldest stored key k_1 and measure the relative reconstruction error of its decayed target value:

$$\frac{\|k_1^\top A_N - \lambda^{N-1} \eta v_1\|_2}{\|\lambda^{N-1} \eta v_1\|_2}.$$

We define effective capacity as the first number of writes N for which this relative error exceeds 1.0. This is a strict oldest-item criterion: capacity fails when interference and forgetting are as large as the remaining signal.

Tensor Cache: Eviction-conditioned Associative Memory for Transformers

Gap	Answer accuracy ($\uparrow$)					Streaming state (kB/seq, $\downarrow$)			
	Full KV	Window KV	StreamingLLM	Infini	TC	Full KV	Infini	TC	TC save
24	1.000 $\pm$ 0.000	0.063 $\pm$ 0.001	0.062 $\pm$ 0.001	0.396 $\pm$ 0.143	1.000 $\pm$ 0.000	398	114	112	71.9%
36	1.000 $\pm$ 0.000	0.063 $\pm$ 0.001	0.063 $\pm$ 0.001	0.206 $\pm$ 0.040	1.000 $\pm$ 0.000	542	114	112	79.3%
48	1.000 $\pm$ 0.000	0.063 $\pm$ 0.001	0.063 $\pm$ 0.001	0.119 $\pm$ 0.039	1.000 $\pm$ 0.000	686	114	112	83.7%

Table 1. Matched-gap synthetic associative recall. Tensor Cache matches Full KV accuracy while using substantially less retained inference state, and outperforms InfiniAttention at nearly identical compact state size.

(a) OpenWebText (130M params, 50K iters, block size 1024, $W = 512$)

Method	NLL at context length L						Peak Mem at 32K [GB]
	1 024	2 048	4 096	8 192	16 384	32 768	
Full KV	4.05	4.46	4.49	5.13	5.40	6.00	1.94
Window KV	4.16	3.98	4.86	4.78	5.08	5.44	0.75
StreamingLLM	4.11	4.66	4.74	4.85	5.35	5.65	0.76
Infini-attention	4.14	4.05	4.49	4.57	5.01	5.42	0.76
Tensor Cache (ours)	3.93	4.00	4.14	4.29	4.55	5.14	0.80

(b) Shakespeare (1M params, 3K iters, block size 256, $W = 128$; mean $\pm$ std over 4 eval seeds)

Method	NLL at context length L						
	256	512	1 024	2 048	4 096	8 192	16 384
Full KV	4.91 $\pm$ 0.57	6.10 $\pm$ 0.44	5.24 $\pm$ 0.57	6.02 $\pm$ 0.79	6.09 $\pm$ 1.05	6.13 $\pm$ 0.82	5.73 $\pm$ 0.30
Window KV	4.91 $\pm$ 0.57	6.25 $\pm$ 0.39	5.48 $\pm$ 0.61	6.16 $\pm$ 0.80	6.27 $\pm$ 1.08	6.24 $\pm$ 0.80	5.82 $\pm$ 0.29
StreamingLLM	4.91 $\pm$ 0.57	6.25 $\pm$ 0.39	5.47 $\pm$ 0.60	6.16 $\pm$ 0.80	6.24 $\pm$ 1.06	6.21 $\pm$ 0.81	5.79 $\pm$ 0.29
Infini-attention	4.80 $\pm$ 0.53	6.06 $\pm$ 0.37	5.14 $\pm$ 0.57	5.79 $\pm$ 0.83	5.87 $\pm$ 1.11	5.92 $\pm$ 0.86	5.51 $\pm$ 0.31
Tensor Cache (ours)	4.79$\pm$0.50	5.98$\pm$0.44	4.96$\pm$0.58	5.60$\pm$0.67	5.66$\pm$1.00	5.70$\pm$0.76	5.32$\pm$0.29

Table 2. Long-context language modeling on (a) OpenWebText and (b) Shakespeare. Streaming NLL on the validation split as a function of evaluation context length; for OpenWebText we additionally report peak GPU memory at $L = 32,768$. All methods within each panel share the same backbone, training data, optimizer, and budget; only the inference-time memory mechanism differs. (a) OpenWebText (130M params): Tensor Cache numbers at $L \leq 4096$ are mean $\pm$ std over 4 eval seeds, longer contexts and all baselines are single-seed. TC attains the lowest mean NLL at every context length except $L = 2048$, where it is statistically indistinguishable from Infini-attention. (b) Shakespeare (1M params): in a converged-regime small-model setting with all methods and contexts evaluated at 4 eval seeds (mean $\pm$ std reported), TC attains the lowest mean NLL at every context length we test, with margins of 0.01–0.22 NLL over the next-best baseline (Infini-attention) and 0.12–0.61 NLL over Window-KV; the margin grows steadily from $L = 512$ to $L = 8192$ and stabilizes at $L = 16,384$ ($64\times$ the trained block size).

We evaluate three regimes. In the *orthonormal-prefix* regime, keys are orthonormal up to the head dimension and then random thereafter, serving as a sanity check. In the *random* regime, all keys are random unit vectors with no decay. In the *decayed* regime, we use the default Tensor Cache write parameters, $\lambda = 0.995$ and $\eta = 0.05$. Results are averaged over five random seeds.

Table 3 shows that raw associative capacity increases with head dimension, as expected for an outer-product memory. The orthonormal-prefix setting remains stable past D writes because the first D keys are orthogonal, while the random setting exposes crosstalk from non-orthogonal keys. Under the decayed setting, capacity is lower because the oldest signal shrinks while newer writes continue to add interference. This supports the interpretation of Tensor Cache as a bounded associative memory: capacity grows with head dimension, number of heads, number of layers, and number

Head dim D	Initialization strategy		
	Ortho-prefix	Random	Decayed
16	31.8 $\pm$ 5.0	19.8 $\pm$ 4.7	15.4 $\pm$ 8.9
32	63.8 $\pm$ 12.0	30.8 $\pm$ 2.6	32.2 $\pm$ 2.9
64	128.4 $\pm$ 18.9	79.0 $\pm$ 18.2	51.8 $\pm$ 6.5
128	240.8 $\pm$ 15.5	143.8 $\pm$ 27.0	84.2 $\pm$ 6.0

Table 3. Raw capacity of a single fast-weight memory head $A \in \mathbb{R}^{D \times D}$, measured as the number of sequential writes until oldest-item relative reconstruction error exceeds 1.0. The decayed setting uses $\lambda = 0.995$, $\eta = 0.05$.

of slots, but retrieval eventually degrades as more unrelated associations are compressed into the same fixed-size state.

Chunk size	Outer rule		Delta rule	
	Acc. $\uparrow$	Time (s) $\downarrow$	Acc. $\uparrow$	Time (s) $\downarrow$
1	1.000	1 145	1.000	1 607
4	1.000	359	1.000	428
8	1.000	198	1.000	243
16	1.000	119	1.000	142
32	1.000	82	1.000	92
64 [†]	0.265	35	0.265	35

[†]Chunk-mean baseline.

Table 4. Chunk-size ablation on synthetic associative recall (matched gap $F = 24$, $W = 12$). The parallel weighted-sum scan of Section 3 (Eq. (8)) preserves perfect retrieval up to chunk size 32 for both update rules; the chunk-mean baseline collapses at 64, where complete episodes fit within a single chunk. Wall-clock times are 1 500-step training on a single RTX 4090.

4.6. Ablations

We ablate the training chunk size and the chunked-mean training shortcut below; the wedge-rule write ablation (negative result) is in Appendix B.4.

Training chunk size and the chunked-mean shortcut.

Table 4 ablates the training chunk size with both update rules using the parallel weighted-sum scan of Section 3 (Eq. (8)). The scan preserves perfect synthetic-recall accuracy at all chunk sizes from $C = 1$ up to $C = 32$ for both rules, with training time decreasing roughly linearly in C (e.g., $14\times$ speedup at $C = 32$ vs. $C = 1$ for the outer rule). At $C = 64$ the chunked-mean shortcut $\bar{k} \otimes \bar{v}$ collapses to 0.265 accuracy on this benchmark, where complete “store–filler–query” episodes fit within a single chunk and the cross-token cross-products dominate the within-chunk write. The outer rule is consistently $\sim 10\text{--}30\%$ faster than the delta rule because it requires no per-token residual computation; both reach the same matched-gap accuracy ceiling. We use the outer rule with $C = 32$ in all main results.

5. Discussion and Limitations

Tensor Cache pairs exact softmax attention over a recent KV window (L1) with a fixed-size outer-product memory A over evicted pairs (L2). Reads use the linear-attention identity $q_t A \approx \sum_i \alpha_i \langle q_t, k_i \rangle v_i$, so the same alignment signal that drives local attention also addresses the compressed past. TC is not a lossless replacement for full attention: Full KV is the exact-content baseline when affordable, and beyond the trained context it is also a positionally-extrapolated baseline whose quality depends on how well RoPE generalizes (Su et al., 2021; Chen et al., 2023; Peng et al., 2024).

Train/inference write rule. We initially used the chunk-mean shortcut $A \leftarrow \lambda A + \eta(\bar{k} \otimes \bar{v})$, which differs from the per-token sum by $C^2 - C$ cross-token outer products

that never appear at inference. The parallel weighted-sum scan of Section 3.4 closes this gap exactly for the outer rule (relative error $< 10^{-7}$). The delta rule is not exactly parallelizable; we use the standard chunk-start residual approximation (Beck et al., 2024), which is empirically tight at moderate chunk sizes (Section 4.6). Even with the parallel scan, an episode-spanning chunk size remains a real failure mode: when an entire “store–filler–query” interval fits within one chunk, intra-chunk reads do not see the corresponding store. This guides chunk-size selection in practice.

Capacity, calibration, and normalization. The matrix A stores associations in superposition, so retrieval degrades as more non-orthogonal associations accumulate. Decay and the optional delta rule mitigate redundancy and stale interference, and we report capacity as a function of head dimension and decay (Section 4.5). TC does not use a read-time normalization vector; mLSTM and Infini-attention do, which bounds read magnitude as the matrix grows but introduces additional confidence-on-OOD failure modes that we observe empirically (calibration paragraph in Section 4.3).

Future directions. Better slot routing could reduce interference between unrelated associations; smaller training chunks or rollout-based training could further tighten the delta-rule chunk-start approximation; FlashLinearAttention-style kernels could reduce read/write overhead; and an optional read-time normalization could be added without changing the eviction-conditioned write semantics. More broadly, TC suggests that KV eviction can be treated as a structured memory operation rather than destructive deletion; richer associative states, denoising filters, or higher-order memories are natural follow-ups.

6. Conclusion

We introduced *Tensor Cache*, a two-level cache that pairs exact softmax attention over a sliding KV window (L1) with a fixed-size outer-product fast-weight memory (L2) fed by KV pairs evicted from the window. Future queries read $q_t A \approx \sum_i \alpha_i \langle q_t, k_i \rangle v_i$, letting older context influence predictions without retaining individual evicted entries. The outer-product memory and the read identity follow Schmidhuber’s fast-weight programmer and the linear-attention reformulation of attention as fast-weight memory; our contribution is to use them as an L2 cache fed exclusively by sliding-window evictions, and to identify and close a chunked-mean training shortcut that introduces $C^2 - C$ spurious cross-token outer products per chunk—closed by a parallel weighted-sum scan equivalent to per-token writes within float32 epsilon. Across systems, real-text long-context language modeling, controlled associative recall, and capacity diagnostics, Tensor Cache provides a practical bounded-state alternative to all existing baselines.

7. Acknowledgments

We would like to thank Manel Baradad and Minyoung Huh for their helpful discussions and advice.

References

- Ba, J., Hinton, G., Mnih, V., Leibo, J. Z., and Ionescu, C. Using fast weights to attend to the recent past. *arXiv preprint arXiv:1610.06258*, 2016. doi: 10.48550/arXiv.1610.06258.
- Beck, M., Pöppel, K., Spanring, M., Auer, A., Prudnikova, O., Kopp, M., Klambauer, G., Brandstetter, J., and Hochreiter, S. xLSTM: Extended long short-term memory. In *Thirty-eighth Conference on Neural Information Processing Systems*, 2024. URL <https://arxiv.org/abs/2405.04517>.
- Beltagy, I., Peters, M. E., and Cohan, A. Longformer: The long-document transformer. *arXiv preprint arXiv:2004.05150*, 2020.
- Bulatov, A., Kuratov, Y., and Burtsev, M. S. Recurrent memory transformer. *arXiv preprint arXiv:2207.06881*, 2022. doi: 10.48550/arXiv.2207.06881.
- Chen, S., Wong, S., Chen, L., and Tian, Y. Extending context window of large language models via positional interpolation. *arXiv preprint arXiv:2306.15595*, 2023.
- Dai, Z., Yang, Z., Yang, Y., Carbonell, J., Le, Q., and Salakhutdinov, R. Transformer-XL: Attentive language models beyond a fixed-length context. *arXiv preprint arXiv:1901.02860*, 2019. doi: 10.48550/arXiv.1901.02860.
- Dao, T. Flashattention-2: Faster attention with better parallelism and work partitioning. *arXiv preprint arXiv:2307.08691*, 2023. doi: 10.48550/arXiv.2307.08691.
- Ding, J., Ma, S., Dong, L., Zhang, X., Huang, S., Wang, W., Zheng, N., and Wei, F. LongNet: Scaling transformers to 1,000,000,000 tokens. *arXiv preprint arXiv:2307.02486*, 2023. doi: 10.48550/arXiv.2307.02486.
- Dong, H., Yang, X., Zhang, Z., Wang, Z., Chi, Y., and Chen, B. Get more with less: Synthesizing recurrence with kv cache compression for efficient llm inference. *arXiv preprint arXiv:2402.09398*, 2024.
- Goel, R., Park, J., Gagrani, M., Jones, D., Morse, M., Langston, H., Lee, M., and Lott, C. CAOTE: KV caching through attention output error based token eviction. *arXiv preprint arXiv:2504.14051*, 2025. doi: 10.48550/arXiv.2504.14051.
- Gokaslan, A. and Cohen, V. Openwebtext corpus. <https://Skylion007.github.io/OpenWebTextCorpus>, 2019.
- Gu, A. and Dao, T. Linear-time sequence modeling with selective state spaces. *arXiv preprint arXiv:2312.00752*, 2023. doi: 10.48550/arXiv.2312.00752.
- Katharopoulos, A., Vyas, A., Pappas, N., and Fleuret, F. Transformers are RNNs: Fast autoregressive transformers with linear attention. In *International Conference on Machine Learning*, pp. 5156–5165, 2020.
- Kwon, W., Li, Z., Zhuang, S., Sheng, Y., Zheng, L., Yu, C. H., Gonzalez, J. E., Zhang, H., and Stoica, I. Efficient memory management for large language model serving with PagedAttention. *arXiv preprint arXiv:2309.06180*, 2023. doi: 10.48550/arXiv.2309.06180.
- Li, Y., Huang, Y., Yang, B., Venkitesh, B., Locatelli, A., Ye, H., Cai, T., Lewis, P., and Chen, D. SnapKV: LLM knows what you are looking for before generation. *arXiv preprint arXiv:2404.14469*, 2024. doi: 10.48550/arXiv.2404.14469.
- Liu, H., Zaharia, M., and Abbeel, P. Ring attention with blockwise transformers for near-infinite context. *arXiv preprint arXiv:2310.01889*, 2023a. doi: 10.48550/arXiv.2310.01889.
- Liu, Y., Li, H., Cheng, Y., Ray, S., Huang, Y., Zhang, Q., Du, K., Yao, J., Lu, S., Ananthanarayanan, G., Maire, M., Hoffmann, H., Holtzman, A., and Jiang, J. CacheGen: KV cache compression and streaming for fast large language model serving. *arXiv preprint arXiv:2310.07240*, 2023b. doi: 10.48550/arXiv.2310.07240.
- Munkhdalai, T., Faruqui, M., and Gopal, S. Leave no context behind: Efficient infinite context transformers with infini-attention. *arXiv preprint arXiv:2404.07143*, 2024. doi: 10.48550/arXiv.2404.07143.
- Peng, B., Quesnelle, J., Fan, H., and Shippole, E. YaRN: Efficient context window extension of large language models. In *International Conference on Learning Representations*, 2024.
- Rae, J. W., Potapenko, A., Jayakumar, S. M., and Lillicrap, T. P. Compressive transformers for long-range sequence modelling. *arXiv preprint arXiv:1911.05507*, 2019. doi: 10.48550/arXiv.1911.05507.
- Schlag, I., Irie, K., and Schmidhuber, J. Linear transformers are secretly fast weight programmers. *arXiv preprint arXiv:2102.11174*, 2021. doi: 10.48550/arXiv.2102.11174.

- Schmidhuber, J. Learning to control fast-weight memories: An alternative to dynamic recurrent networks. *Neural Computation*, 4(1):131–139, 1992.
- Shazeer, N. Fast transformer decoding: One write-head is all you need. *arXiv preprint arXiv:1911.02150*, 2019.
- Su, J., Lu, Y., Pan, S., Murtadha, A., Wen, B., and Liu, Y. Roformer: Enhanced transformer with rotary position embedding. *arXiv preprint arXiv:2104.09864*, 2021.
- Sun, Y., Dong, L., Huang, S., Ma, S., Xia, Y., Xue, J., Wang, J., and Wei, F. Retentive network: A successor to transformer for large language models. *arXiv preprint arXiv:2307.08621*, 2023. doi: 10.48550/arXiv.2307.08621.
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, L., and Polosukhin, I. Attention is all you need. In *Advances in Neural Information Processing Systems*, 2017.
- Widrow, B. and Hoff, M. E. Adaptive switching circuits. In *1960 IRE WESCON Convention Record*, volume 4, pp. 96–104. IRE, 1960.
- Wu, Y., Rabe, M. N., Hutchins, D., and Szegedy, C. Memorizing transformers. *arXiv preprint arXiv:2203.08913*, 2022. doi: 10.48550/arXiv.2203.08913.
- Xiao, G., Tian, Y., Chen, B., Han, S., and Lewis, M. Efficient streaming language models with attention sinks. *arXiv preprint arXiv:2309.17453*, 2023. doi: 10.48550/arXiv.2309.17453.
- Yang, S., Wang, B., Zhang, Y., Shen, Y., and Kim, Y. Parallelizing linear transformers with the delta rule over sequence length. In *Advances in Neural Information Processing Systems*, 2024. URL <https://arxiv.org/abs/2406.06484>.
- Zhang, Z., Sheng, Y., Zhou, T., Chen, T., Zheng, L., Cai, R., Song, Z., Tian, Y., Ré, C., Barrett, C., Wang, Z., and Chen, B. H₂O: Heavy-hitter oracle for efficient generative inference of large language models. *arXiv preprint arXiv:2306.14048*, 2023. doi: 10.48550/arXiv.2306.14048.

A. Architecture Extensions

This appendix collects the optional Tensor Cache variants deferred from Section 3.

A.1. Multi-slot memory and routing

In the multi-slot setting, each layer maintains S memory matrices per head:

$$A \in \mathbb{R}^{B \times H \times S \times D \times D}. \quad (10)$$

Each head has S learned slot keys $K^{\text{slot}} \in \mathbb{R}^{H \times S \times D}$. A query produces router logits and Top- k softmax weights $\pi_{t,s}$:

$$\ell_{t,s} = \frac{\langle \hat{q}_t, \hat{K}_s^{\text{slot}} \rangle}{\sqrt{D} \tau}, \quad \pi_{t,s} = \text{TopKSoftmax}(\ell_{t,\cdot})_s, \quad (11)$$

where $\hat{\cdot}$ denotes L2-normalization and τ is a temperature. The read mixes the per-slot reads:

$$r_t = \sum_{s=1}^S \pi_{t,s} (q_t A_s). \quad (12)$$

Writes update a single routed slot using either the write key k_w or the current query q_t , implemented with straight-through Gumbel-Softmax for differentiability through the routing decision.

A.2. Two-timescale memory

Each layer optionally maintains a fast and slow memory ($A_{\text{fast}}, A_{\text{slow}}$) with separate per-head decay/lr parameters and reads

$$r_t = r_t^{\text{fast}} + \alpha r_t^{\text{slow}}, \quad \alpha = \sigma(\alpha_{\text{logit}}). \quad (13)$$

This separates short-horizon binding from long-horizon retention without changing the read or fusion path. Both timescales receive the same write source on eviction.

A.3. Layer selection and stability knobs

Tensor Cache can be applied to all layers, no layers (effectively reducing to Window KV with the gate forced off), or only the top- k layers. We additionally support optional query L2-normalization, key L2-normalization, and a value scaling factor; these are exposed as ablation knobs and disabled by default in the main results.

B. Additional Experiments

This appendix provides extended synthetic recall results, stress tests, and training efficiency measurements that complement the main paper.

B.1. Long-Gap Synthetic Associative Recall

To probe retention at substantially larger delays, we run a long-gap synthetic associative-recall benchmark with gaps $\{128, 256, 512, 1024\}$. This setting uses a simplified single-episode, single-store task for tractability and compares Full KV, InfiniAttention, and Tensor Cache. Through gaps 128–512, all three methods achieve perfect answer accuracy, while Tensor Cache remains fixed at 112 kB per sequence and Full KV grows from 286 kB to 1054 kB. At gap 1024, InfiniAttention begins to degrade (0.872 ± 0.106), while Tensor Cache retains 0.959 ± 0.058 accuracy at 94.6% state savings relative to Full KV.

Tensor Cache: Eviction-conditioned Associative Memory for Transformers

Gap	Answer accuracy ($\uparrow$)			Streaming state (kB/seq, $\downarrow$)			TC save
	Full KV	Infini	TC	Full KV	Infini	TC	
128	1.000 ± 0.000	1.000 ± 0.000	1.000 ± 0.000	286	114	112	60.8%
256	1.000 ± 0.000	1.000 ± 0.000	1.000 ± 0.000	542	114	112	79.3%
512	1.000 ± 0.000	1.000 ± 0.000	1.000 ± 0.000	1054	114	112	89.4%
1024	1.000 ± 0.000	0.872 ± 0.106	0.959 ± 0.058	2078	114	112	94.6%

Table 5. Long-gap synthetic associative recall in the simplified single-store setting. Values are mean $\pm$ standard deviation over 3 seeds. Tensor Cache retains near-perfect accuracy through gap 1024 while maintaining constant 112 kB per-sequence state.

B.2. Synthetic Stress Tests

We evaluate a multistore synthetic variant to test robustness under increased memory load. In the multistore stress test at gap 24 with two stored pairs per episode, Tensor Cache achieves 0.563 ± 0.003 answer accuracy, compared with 0.558 ± 0.004 for Full KV, 0.464 ± 0.118 for InfiniAttention, and near-chance performance for strictly local baselines. Because Full KV is not fully saturated in this harder configuration, we treat this result as a robustness check rather than a main comparison.

Case	Answer accuracy ($\uparrow$)					State (kB/seq, $\downarrow$)			TC save
	Full	Window	SLLM	Infini	TC	Full	Infini	TC	
Multistore	0.558 ± 0.004	0.125 ± 0.005	0.128 ± 0.003	0.464 ± 0.118	0.563 ± 0.003	434	114	112	74.2%

Table 6. Multistore synthetic stress test at gap 24 with two stored pairs per episode. SLLM denotes StreamingLLM. Tensor Cache matches Full KV (within seed variance) while using 74.2% less retained inference state.

B.3. Training Efficiency

Tables 7 and 8 report training time, training memory, evaluation memory, and per-sequence streaming state for the synthetic benchmarks. Tensor Cache trains faster than InfiniAttention at all gaps, but uses more training memory due to the associative-memory state maintained during backpropagation. At evaluation time, both methods maintain nearly identical compact streaming state (112–114 kB per sequence).

Gap	Method	Training		Inference	
		Time (s) $\downarrow$	Mem (GB) $\downarrow$	Mem (GB) $\downarrow$	State (kB) $\downarrow$
24	Full KV	6.8 ± 0.0	0.162	0.046	398
	Infini	$1\,007.7 \pm 4.2$	0.608	0.035	114
	TC	798.4 ± 1.1	1.374	0.037	112
36	Full KV	7.3 ± 0.2	0.216	0.052	542
	Infini	$1\,372.3 \pm 3.5$	0.825	0.035	114
	TC	$1\,102.5 \pm 3.0$	1.867	0.038	112
48	Full KV	7.1 ± 0.1	0.257	0.057	686
	Infini	$1\,754.4 \pm 2.8$	1.025	0.036	114
	TC	$1\,410.1 \pm 5.4$	2.345	0.038	112

Table 7. Synthetic benchmark efficiency at matched gaps. TC trains faster than Infini-attention but uses more training memory; both maintain constant-size inference state. Train time: mean $\pm$ std over 3 seeds; remaining columns are deterministic (std = 0).

Tensor Cache: Eviction-conditioned Associative Memory for Transformers

Gap	Method	Training		Inference	
		Time (s) ↓	Mem (GB) ↓	Mem (GB) ↓	State (kB) ↓
128	Full KV	57.9 ± 3.0	0.124	0.044	286
	Infini	4 191.3 ± 206.5	0.445	0.035	114
	TC	2 956.6 ± 276.9	0.996	0.037	112
256	Full KV	53.5 ± 0.2	0.216	0.052	542
	Infini	7 574.7 ± 6.5	0.825	0.035	114
	TC	5 056.9 ± 3.9	1.867	0.038	112
512	Full KV	54.4 ± 0.4	0.381	0.072	1 054
	Infini	15 218.3 ± 127.7	1.561	0.036	114
	TC	9 993.0 ± 160.9	3.588	0.039	112
1024	Full KV	62.7 ± 0.7	0.725	0.114	2 078
	Infini	30 621.6 ± 208.9	3.054	0.039	114
	TC	20 299.1 ± 91.4	7.050	0.041	112

Table 8. Long-gap synthetic training efficiency. TC trains faster than Infini-attention but uses more training memory; both maintain nearly identical compact streaming state at inference. Train time: mean ± std over 3 seeds; remaining columns are deterministic (std = 0).

B.4. Write-Rule Ablation: Wedge Product (Negative Result)

A natural geometric question is whether the rank-1 outer-product write $A \leftarrow \lambda A + \eta k v^\top$ should be replaced by the antisymmetric *wedge* (exterior-algebra) write

$$A \leftarrow \lambda A + \eta (k v^\top - v k^\top), \quad (14)$$

on the intuition that the wedge product is directional ($k \rightarrow v$ but not $v \rightarrow k$) and stores only one independent triangle of A . We implement the wedge variant in both the streaming per-token path and the parallel weighted-sum scan; the chunked scan agrees with per-token streaming within float32 epsilon, and A is exactly antisymmetric throughout training ($\|A + A^\top\|_\infty = 0$).

The argument breaks once the read formula is examined. Because TC reads with $q^\top A$, the plain outer rule already gives directional retrieval $q^\top (k v^\top) = \langle q, k \rangle v$ in the read direction we use; the $v \rightarrow k$ path corresponds to the right-multiply $A q$ that TC never performs. The wedge replaces this with $q^\top (k v^\top - v k^\top) = \langle q, k \rangle v - \langle q, v \rangle k$, introducing a new $-\langle q, v \rangle k$ crosstalk term that is non-negligible whenever queries are partially aligned with stored values. Since k and v are linear projections of the same hidden state in attention, this alignment is the typical case rather than an edge case.

Empirically, on the matched-gap synthetic associative-recall benchmark of Section 4.3 ($g = 24$, $W = 12$, 3 seeds, 300 training steps, identical configuration in all other respects), the wedge rule is consistently worse than both the outer and delta rules on accuracy, NLL, and seed-to-seed stability (Table 9). Mean answer accuracy is 0.998 for outer and 1.000 for delta, versus 0.979 for wedge; mean NLL is 0.021 for outer and 0.001 for delta, versus 0.150 for wedge—roughly $7\times$ worse than outer and $150\times$ worse than delta. Two of three wedge seeds reach ≥ 0.999 accuracy, indicating it is capable of solving the task; the third seed converges substantially slower than under either of the other rules at the same compute budget. Per-step training time is $\sim 30\%$ higher for both wedge and delta than for outer (wedge from doubling the per-token outer products, delta from the residual matmul $v - k^\top A$). We retain the outer rule in all main results—it matches delta on the matched-gap task once chunk size $C = 32$ (Table 4) while being cheapest—and we keep the wedge available behind `tm.update_rule="wedge"` as a negative-result ablation.

Write rule	Acc. ↑	NLL ↓	Train time (s) ↓	Per-seed acc.
Outer	0.998 ± 0.003	0.021 ± 0.020	201.0 ± 1.4	0.994/1.000/0.999
Delta	1.000 ± 0.000	0.001 ± 0.000	264.8 ± 1.9	1.000/1.000/1.000
Wedge	0.979 ± 0.029	0.150 ± 0.156	260.5 ± 1.7	0.938/0.999/1.000

Table 9. Write-rule ablation on the matched-gap synthetic associative-recall task ($g = 24$, $W = 12$, 3 seeds, 300 training steps, all other knobs identical). The wedge rule retains the antisymmetric structure $A = -A^\top$ throughout training, but introduces a $-\langle q, v \rangle k$ crosstalk term in the readout that the outer and delta rules do not have. Across seeds, wedge is less accurate and substantially less stable than both alternatives, while costing the same training time as delta. Outer and delta both solve the task perfectly within seed noise; the chunk-size ablation in Table 4 confirms they remain matched at $C = 32$, which is why we use the cheaper outer rule in all main results.